

In this problem of multiclass classification, we are going to build a neural network to classify images of different items of clothing.

```
import tensorflow as tf
from tensorflow.keras.datasets import fashion_mnist

#The data has already been sorted into training and test sets for us.
(train_data, train_labels), (test_data, test_labels) = fashion_mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
29515/29515 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
26421880/26421880 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
5148/5148 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4422102/4422102 [=====] - 0s 0us/step

#Showing the first training example
print(f"Training Sample:\n{train_data[0]}\n")
print(f"Training Label:\n{train_labels[0]}\n")
```

Training Sample:

```
[ [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
    0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      0 1 4 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      54 0 0 0 1 3 4 0 0 0 3 0 0 0 0 0 0 0
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      144 123 23 0 0 0 0 12 10 0 0 0 0 0 0 0 0
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      107 156 161 109 64 23 77 130 72 15 0 0 0 0 0 0
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      216 163 127 121 122 146 141 88 172 66 0 0 0 0 0 0
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      223 223 215 213 164 127 123 196 229 0 1 1 1 0 200 232 232 233 229
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      235 227 224 222 224 221 223 245 173 0 0 0 0 0 0 183 225 216 223 228
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      180 212 210 211 213 223 220 243 202 0 0 0 0 0 0 193 228 218 213 198
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      169 227 208 218 224 212 226 197 209 52 1 3 0 12 219 220 212 218 192
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      198 221 215 213 222 220 245 119 167 56 6 0 99 244 222 220 218 203
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      232 213 218 223 234 217 217 209 92 0 4 0 0 55 236 228 230 228 240
    [ 0 0 1 4 6 7 2 0 0 0 0 0 0 0 0 0 0 0
      222 221 216 223 229 215 218 255 77 0 0 0 0 237 226 217 223 222 219
    [ 0 3 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      211 218 224 223 219 215 224 244 159 0 62 145 204 228 207 213 221 218 208
    [ 0 0 0 0 18 44 82 107 189 228 220 222 217 226 200 205 211 230
      224 234 176 188 250 248 233 238 215 0 204 214 208 209 200 159 245 193 206 223
    [ 0 57 187 208 224 221 224 208 204 214 208 209 200 159 245 193 206 223
      255 255 221 234 221 211 220 232 246 0 205 205 205 220 240 80 150 255 229 221
    [ 3 202 228 224 221 211 211 214 205 205 205 220 240 80 150 255 229 221
      188 154 191 210 204 209 222 228 225 0 220 194 215 217 241 65 73 106 117
    [ 98 233 198 210 222 229 229 234 249 220 194 215 217 241 65 73 106 117
      168 219 221 215 217 223 223 224 229 29 185 197 206 198 213 240 195 227 245
    [ 75 204 212 204 193 205 211 225 216 185 197 206 198 213 240 195 227 245
      239 223 218 212 209 222 220 221 230 67 194 192 202 214 219 221 220 236 225 216
    [ 48 203 183 194 213 197 185 190 194 192 202 214 219 221 220 236 225 216
      199 206 186 181 177 172 181 205 206 115 210 213 207 211 210 200 196 194 191
    [ 0 122 219 193 179 171 183 196 204 210 213 207 211 210 200 196 194 191
      195 191 198 192 176 156 167 177 210 92 181 185 188 189 188 193 198 204 209
    [ 0 0 74 189 212 191 175 172 175 181 185 188 189 188 193 198 204 209
      210 210 211 188 188 194 192 216 170 0 0 0 0 0 0 0 0 0 0
    [ 2 0 0 0 66 200 222 237 239 242 246 243 244 221 220 193 191 179
      182 182 181 176 166 168 99 58 0 0 61 44 72 41 35 0 0 0 0 0 0
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
    [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
      0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
```

Checking the shapes of the test and train data

```
train_data.shape, train_labels.shape
```

```
((60000, 28, 28), (60000,))
```


```
test_data.shape, test_labels.shape
```

```
((10000, 28, 28), (10000,))
```

```
#Plotting a single sample
```

```
import matplotlib.pyplot as plt
```

```
plt.imshow(train_data[0]);
```



```
#checking its label
```

```
train_labels[0]
```

```
9
```

```
#creating human readable labels of given training data
```

```
class_names = ["T-shirt/top", "Trousers", "Pullover", "Dress", "Coat", "Sandal", "Shirt", "Sneaker", "Bag", "Ankle boot"]
```

```
len(class_names)
```

```
10
```

```
#Plotting an example image and its label
```

```
index = 18
```

```
plt.imshow(train_data[index], cmap=plt.cm.binary)
```

```
plt.title(class_names[train_labels[index]])
```

```
Text(0.5, 1.0, 'Shirt')
```



```
import random
plt.figure(figsize=(7,7))
for i in range(4):
    ax = plt.subplot(2, 2, i+1)
    random_index = random.choice(range(len(train_data)))
    plt.imshow(train_data[random_index], cmap=plt.cm.binary)
    plt.title(class_names[train_labels[random_index]])
    plt.axis(False)
```



▼ Now lets build the multiclassification model

```
#setting the random seed
tf.random.set_seed(42)

#create the model
model_1 = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28)),
    tf.keras.layers.Dense(4, activation="relu"),
    tf.keras.layers.Dense(4, activation="relu"),
    tf.keras.layers.Dense(10, activation="softmax")
])

#compiling the model
model_1.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                optimizer=tf.keras.optimizers.Adam(),
                metrics = ['accuracy'])

#Fitting the model
history = model_1.fit(train_data,
                      train_labels,
                      epochs=10,
                      validation_data=(test_data, test_labels))
```

```
Epoch 1/10
1875/1875 [=====] - 10s 4ms/step - loss: 2.1452 - accuracy: 0.1860 - val_loss: 1.7850 - val_
Epoch 2/10
1875/1875 [=====] - 7s 4ms/step - loss: 1.7162 - accuracy: 0.2665 - val_loss: 1.6408 - val_a
Epoch 3/10
1875/1875 [=====] - 4s 2ms/step - loss: 1.6369 - accuracy: 0.2772 - val_loss: 1.5692 - val_a
Epoch 4/10
1875/1875 [=====] - 4s 2ms/step - loss: 1.5362 - accuracy: 0.3207 - val_loss: 1.4878 - val_a
Epoch 5/10
```

```

1875/1875 [=====] - 5s 2ms/step - loss: 1.4750 - accuracy: 0.3490 - val_loss: 1.4376 - val_a
Epoch 6/10
1875/1875 [=====] - 3s 2ms/step - loss: 1.4402 - accuracy: 0.3675 - val_loss: 1.4502 - val_a
Epoch 7/10
1875/1875 [=====] - 4s 2ms/step - loss: 1.4186 - accuracy: 0.3839 - val_loss: 1.5016 - val_a
Epoch 8/10
1875/1875 [=====] - 5s 3ms/step - loss: 1.4182 - accuracy: 0.3844 - val_loss: 1.4111 - val_a
Epoch 9/10
1875/1875 [=====] - 4s 2ms/step - loss: 1.4217 - accuracy: 0.3868 - val_loss: 1.4088 - val_a
Epoch 10/10
1875/1875 [=====] - 3s 2ms/step - loss: 1.3855 - accuracy: 0.3961 - val_loss: 1.3860 - val_a

```

#checking the model summary

```
model_1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 4)	3140
dense_1 (Dense)	(None, 4)	20
dense_2 (Dense)	(None, 10)	50
Total params: 3,210		
Trainable params: 3,210		
Non-trainable params: 0		

▼ Now lets try to improve the accuracy by standardising or normalizing the data (between 0 and 1)

```
#checking the min and max values of the training data
train_data.min(),train_data.max()
```

```
(0, 255)
```

```
#normalizing the training and testing data
norm_train_data = train_data/255.0
norm_test_data = test_data/255.0
```

```
#checking our normalized data
norm_train_data.max(), norm_test_data.min()
```

```
(1.0, 0.0)
```

#Using this normalised data on the same model we built above

```
#setting random seed
tf.random.set_seed(42)
```

```
#building the model
model_2 = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape = (28,28)),
    tf.keras.layers.Dense(4, activation = "relu"),
    tf.keras.layers.Dense(4, activation = "relu"),
    tf.keras.layers.Dense(10, activation = "softmax")
])
```

```
#compiling the model
model_2.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                optimizer=tf.keras.optimizers.Adam(),
                metrics = ['accuracy'])
```

```
#fitting the model
norm_history = model_2.fit(norm_train_data,
                           train_labels,
                           epochs=10,
                           validation_data = (norm_test_data, test_labels))
```

```

Epoch 1/10
1875/1875 [=====] - 6s 3ms/step - loss: 1.6080 - accuracy: 0.4029 - val_loss: 1.1924 - val_a
Epoch 2/10
1875/1875 [=====] - 4s 2ms/step - loss: 1.0581 - accuracy: 0.6031 - val_loss: 0.9879 - val_a
Epoch 3/10

```

```

1875/1875 [=====] - 4s 2ms/step - loss: 0.8989 - accuracy: 0.6571 - val_loss: 0.8808 - val_a
Epoch 4/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.8250 - accuracy: 0.6763 - val_loss: 0.8377 - val_a
Epoch 5/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.7917 - accuracy: 0.6867 - val_loss: 0.8135 - val_a
Epoch 6/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.7653 - accuracy: 0.7173 - val_loss: 0.7885 - val_a
Epoch 7/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.7322 - accuracy: 0.7395 - val_loss: 0.7580 - val_a
Epoch 8/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.7015 - accuracy: 0.7476 - val_loss: 0.7343 - val_a
Epoch 9/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.6805 - accuracy: 0.7575 - val_loss: 0.7141 - val_a
Epoch 10/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.6637 - accuracy: 0.7624 - val_loss: 0.7001 - val_a

```



▼ Plotting the loss curves for normalized data and non-normalized data

```
import pandas as pd
```

```
#Plotting the non-normalized data loss curve
pd.DataFrame(history.history).plot(title = "Non_Normalized")
```

```
#Plotting the normalized data loss curve
pd.DataFrame(history.history).plot(title = "Normalized")
```

```
<Axes: title={'center': 'Normalized'}>
```



▼ Finding the ideal learning rate using the callback

```

#Using this normalised data on the same model we built above but with callback method

#setting random seed
tf.random.set_seed(42)

#building the model
model_3 = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape = (28,28)),
    tf.keras.layers.Dense(4, activation = "relu"),
    tf.keras.layers.Dense(4, activation = "relu"),
    tf.keras.layers.Dense(10, activation = "softmax")
])

#compiling the model
model_3.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                optimizer=tf.keras.optimizers.Adam(),
                metrics = ['accuracy'])

#creaaating the learning rate callback
lr_scheduler = tf.keras.callbacks.LearningRateScheduler(lambda epoch : 1e-3 * 10**(epoch/20))

#fitting the model
lr_norm_history = model_3.fit(norm_train_data,
                              train_labels,
                              epochs=40,
                              validation_data = (norm_test_data, test_labels),
                              callbacks=[lr_scheduler])

Epoch 12/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5546 - accuracy: 0.7938 - val_loss: 0.6100 - val_a
Epoch 13/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5538 - accuracy: 0.7939 - val_loss: 0.6563 - val_a
Epoch 14/40
1875/1875 [=====] - 5s 3ms/step - loss: 0.5568 - accuracy: 0.7910 - val_loss: 0.5970 - val_a
Epoch 15/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5566 - accuracy: 0.7927 - val_loss: 0.5727 - val_a
Epoch 16/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5603 - accuracy: 0.7911 - val_loss: 0.5771 - val_a
Epoch 17/40
1875/1875 [=====] - 5s 3ms/step - loss: 0.5584 - accuracy: 0.7926 - val_loss: 0.6058 - v
Epoch 18/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5637 - accuracy: 0.7922 - val_loss: 0.5792 - v
Epoch 19/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5637 - accuracy: 0.7907 - val_loss: 0.5902 - v
Epoch 20/40
1875/1875 [=====] - 5s 2ms/step - loss: 0.5663 - accuracy: 0.7909 - val_loss: 0.5899 - v
Epoch 21/40
1875/1875 [=====] - 5s 2ms/step - loss: 0.5722 - accuracy: 0.7885 - val_loss: 0.6066 - val_a
Epoch 22/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5775 - accuracy: 0.7864 - val_loss: 0.5940 - val_a
Epoch 23/40
1875/1875 [=====] - 5s 3ms/step - loss: 0.5867 - accuracy: 0.7831 - val_loss: 0.6235 - val_a
Epoch 24/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5864 - accuracy: 0.7849 - val_loss: 0.6453 - val_a
Epoch 25/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5939 - accuracy: 0.7829 - val_loss: 0.6242 - val_a
Epoch 26/40
1875/1875 [=====] - 6s 3ms/step - loss: 0.6108 - accuracy: 0.7779 - val_loss: 0.8094 - val_a
Epoch 27/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.6161 - accuracy: 0.7770 - val_loss: 0.6748 - val_a
Epoch 28/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.6295 - accuracy: 0.7722 - val_loss: 0.6362 - val_a
Epoch 29/40

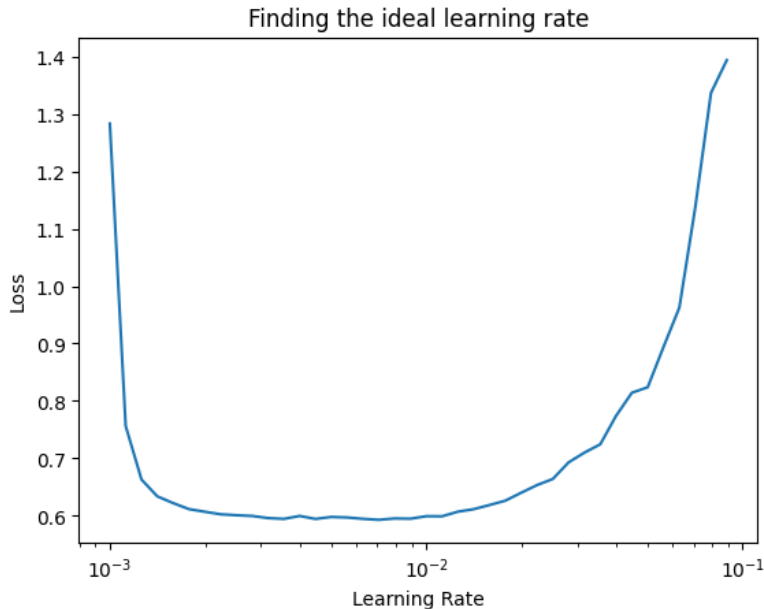
```

```
1875/1875 [=====] - 4s 2ms/step - loss: 1.3824 - accuracy: 0.3855 - val_loss: 1.4553 - val_a
```

```
#plotting the learning rate decay curve
import numpy as np
import matplotlib.pyplot as plt

lrs = 1e-3 * (10**(tf.range(40)/20))
plt.semilogx(lrs, lr_norm_history.history["loss"])
plt.xlabel("Learning Rate")
plt.ylabel("Loss")
plt.title("Finding the ideal learning rate")
```

```
Text(0.5, 1.0, 'Finding the ideal learning rate')
```



▼ Lets now rebuild our model with the ideal learning rate

```
#Using this normalised data on the same model we built above but with callback method and ideal learning rate
```

```
#setting random seed
tf.random.set_seed(42)
```

```
#building the model
model_4 = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape = (28,28)),
    tf.keras.layers.Dense(4, activation = "relu"),
    tf.keras.layers.Dense(4, activation = "relu"),
    tf.keras.layers.Dense(10, activation = "softmax")
])
```

```
#compiling the model
model_4.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                optimizer=tf.keras.optimizers.Adam(lr=0.001),
                metrics = ['accuracy'])
```

```
#fitting the model
lr_norm_history = model_4.fit(norm_train_data,
                              train_labels,
                              epochs=40,
                              validation_data = (norm_test_data, test_labels))
```

```
WARNING:absl:lr` is deprecated in Keras optimizer, please use `learning_rate` or use the legacy optimizer, e.g., tf.k
Epoch 1/40
1875/1875 [=====] - 5s 2ms/step - loss: 1.3171 - accuracy: 0.4884 - val_loss: 0.9585 - val_a
Epoch 2/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.8798 - accuracy: 0.6441 - val_loss: 0.8718 - val_a
Epoch 3/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.8215 - accuracy: 0.6564 - val_loss: 0.8271 - val_a
Epoch 4/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.7902 - accuracy: 0.6858 - val_loss: 0.7929 - val_a
Epoch 5/40
1875/1875 [=====] - 5s 3ms/step - loss: 0.7597 - accuracy: 0.7195 - val_loss: 0.7652 - val_a
Epoch 6/40
1875/1875 [=====] - 3s 2ms/step - loss: 0.7404 - accuracy: 0.7331 - val_loss: 0.7708 - val_a
Epoch 7/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.7263 - accuracy: 0.7423 - val_loss: 0.7695 - val_a
Epoch 8/40
```

```

1875/1875 [=====] - 5s 2ms/step - loss: 0.7077 - accuracy: 0.7509 - val_loss: 0.7184 - val_a
Epoch 9/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.6164 - accuracy: 0.7886 - val_loss: 0.6267 - val_a
Epoch 10/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5807 - accuracy: 0.8024 - val_loss: 0.6092 - val_a
Epoch 11/40
1875/1875 [=====] - 5s 2ms/step - loss: 0.5665 - accuracy: 0.8081 - val_loss: 0.6037 - val_a
Epoch 12/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5559 - accuracy: 0.8114 - val_loss: 0.5957 - val_a
Epoch 13/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5495 - accuracy: 0.8120 - val_loss: 0.5907 - val_a
Epoch 14/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5435 - accuracy: 0.8153 - val_loss: 0.5884 - val_a
Epoch 15/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5389 - accuracy: 0.8170 - val_loss: 0.5807 - val_a
Epoch 16/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5353 - accuracy: 0.8172 - val_loss: 0.5780 - val_a
Epoch 17/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5300 - accuracy: 0.8192 - val_loss: 0.5801 - val_a
Epoch 18/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5285 - accuracy: 0.8188 - val_loss: 0.5752 - val_a
Epoch 19/40
1875/1875 [=====] - 3s 2ms/step - loss: 0.5255 - accuracy: 0.8208 - val_loss: 0.5760 - val_a
Epoch 20/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5227 - accuracy: 0.8225 - val_loss: 0.5840 - val_a
Epoch 21/40
1875/1875 [=====] - 5s 3ms/step - loss: 0.5202 - accuracy: 0.8215 - val_loss: 0.5786 - val_a
Epoch 22/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5176 - accuracy: 0.8235 - val_loss: 0.5776 - val_a
Epoch 23/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5165 - accuracy: 0.8232 - val_loss: 0.5770 - val_a
Epoch 24/40
1875/1875 [=====] - 5s 3ms/step - loss: 0.5133 - accuracy: 0.8250 - val_loss: 0.5666 - val_a
Epoch 25/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5112 - accuracy: 0.8249 - val_loss: 0.5824 - val_a
Epoch 26/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5126 - accuracy: 0.8245 - val_loss: 0.5653 - val_a
Epoch 27/40
1875/1875 [=====] - 5s 3ms/step - loss: 0.5085 - accuracy: 0.8258 - val_loss: 0.5614 - val_a
Epoch 28/40
1875/1875 [=====] - 4s 2ms/step - loss: 0.5069 - accuracy: 0.8263 - val_loss: 0.5616 - v
Epoch 29/40

```

▼ Evaluating our multiclass classification model

```

# Note: The following confusion matrix code is a remix of Scikit-Learn's
# plot_confusion_matrix function - https://scikit-learn.org/stable/modules/generated/sklearn.metrics.plot_confusion_matrix
# and Made with ML's introductory notebook - https://github.com/GokuMohandas/MadeWithML/blob/main/notebooks/08_Neural_Netw
import itertools
from sklearn.metrics import confusion_matrix

# Our function needs a different name to sklearn's plot_confusion_matrix
def make_confusion_matrix(y_true, y_pred, classes=None, figsize=(10, 10), text_size=15):
    """Makes a labelled confusion matrix comparing predictions and ground truth labels.

    If classes is passed, confusion matrix will be labelled, if not, integer class values
    will be used.

    Args:
        y_true: Array of truth labels (must be same shape as y_pred).
        y_pred: Array of predicted labels (must be same shape as y_true).
        classes: Array of class labels (e.g. string form). If `None`, integer labels are used.
        figsize: Size of output figure (default=(10, 10)).
        text_size: Size of output figure text (default=15).

    Returns:
        A labelled confusion matrix plot comparing y_true and y_pred.

    Example usage:
        make_confusion_matrix(y_true=test_labels, # ground truth test labels
                              y_pred=y_preds, # predicted labels
                              classes=class_names, # array of class label names
                              figsize=(15, 15),
                              text_size=10)
    """
    # Create the confusion matrix
    cm = confusion_matrix(y_true, y_pred)
    cm_norm = cm.astype("float") / cm.sum(axis=1)[:, np.newaxis] # normalize it
    n_classes = cm.shape[0] # find the number of classes we're dealing with

    # Plot the figure and make it pretty

```



```

fig, ax = plt.subplots(figsize=figsize)
cax = ax.matshow(cm, cmap=plt.cm.Blues) # colors will represent how 'correct' a class is, darker == better
fig.colorbar(cax)

# Are there a list of classes?
if classes:
    labels = classes
else:
    labels = np.arange(cm.shape[0])

# Label the axes
ax.set(title="Confusion Matrix",
        xlabel="Predicted label",
        ylabel="True label",
        xticks=np.arange(n_classes), # create enough axis slots for each class
        yticks=np.arange(n_classes),
        xticklabels=labels, # axes will labeled with class names (if they exist) or ints
        yticklabels=labels)

# Make x-axis labels appear on bottom
ax.xaxis.set_label_position("bottom")
ax.xaxis.tick_bottom()

# Set the threshold for different colors
threshold = (cm.max() + cm.min()) / 2.

# Plot the text on each cell
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, f"{cm[i, j]} ({cm_norm[i, j]*100:.1f}%)",
             horizontalalignment="center",
             color="white" if cm[i, j] > threshold else "black",
             size=text_size)

class_names

['T-shirt/top',
 'Trousers',
 'Pullover',
 'Dress',
 'Coat',
 'Sandal',
 'Shirt',
 'Sneaker',
 'Bag',
 'Ankle boot']

test_labels

array([9, 2, 1, ..., 8, 1, 5], dtype=uint8)


#Lets make predictions with our model
y_probs = model_4.predict(norm_test_data)

#Viewing the first 5 predictions
y_probs[:5]

313/313 [=====] - 1s 2ms/step
array([[1.65595619e-14, 1.33894042e-11, 8.38167106e-16, 1.64987485e-25,
        8.12260745e-11, 3.35326016e-01, 4.41123180e-12, 1.15524903e-01,
        2.26949146e-06, 5.49146891e-01],
       [7.65582536e-06, 3.52378776e-17, 9.56999719e-01, 6.71652911e-09,
        1.16431816e-02, 1.29238295e-11, 3.13492306e-02, 5.53316900e-34,
        1.54084987e-07, 7.98325974e-24],
       [2.39907950e-03, 9.78370845e-01, 7.93917934e-06, 1.19544845e-02,
        2.48201486e-05, 2.62084045e-03, 9.18941732e-05, 4.43833135e-03,
        8.30943973e-05, 8.71042357e-06],
       [2.39907950e-03, 9.78370845e-01, 7.93917934e-06, 1.19544845e-02,
        2.48201486e-05, 2.62084045e-03, 9.18941732e-05, 4.43833135e-03,
        8.30943973e-05, 8.71042357e-06],
       [1.35179564e-01, 9.24505457e-06, 5.09983063e-01, 5.59618790e-03,
        1.30348951e-02, 1.49875405e-05, 3.35897774e-01, 5.20340694e-18,
        2.84260197e-04, 1.92658254e-13]], dtype=float32)

#plotting the first data object
plt.imshow(train_data[0]);

```



```
#Making its prediction
y_probs[0], tf.argmax(y_probs[0]), class_names[tf.argmax(y_probs[0])]

(array([1.6559562e-14, 1.3389404e-11, 8.3816711e-16, 1.6498748e-25,
        8.1226074e-11, 3.3532602e-01, 4.4112318e-12, 1.1552490e-01,
        2.2694915e-06, 5.4914689e-01], dtype=float32),
<tf.Tensor: shape=(), dtype=int64, numpy=9>,
'Ankle boot')
```

```
#Now lets convert all the prediction probabilities into integers
y_preds = y_probs.argmax(axis=1)
```

```
#Viewing the first 1 predictions
y_preds[0], class_names[tf.argmax(y_probs[0])]

(9, 'Ankle boot')
```

```
test_labels

array([9, 2, 1, ..., 8, 1, 5], dtype=uint8)
```

Confusion Matrix

```
from sklearn.metrics import confusion_matrix
confusion_matrix(y_true=test_labels,
                  y_pred=y_preds)

array([[829,  4, 19, 61,  1,  1, 72,  0, 13,  0],
       [ 9, 927,  4, 52,  2,  1,  5,  0,  0,  0],
       [29,  1, 743, 18, 105,  0, 103,  0,  1,  0],
       [61, 15, 31, 785, 50,  1, 51,  0,  6,  0],
       [ 1,  0, 190, 42, 636,  0, 125,  0,  6,  0],
       [ 0,  0,  0,  1,  1, 943,  0, 28,  7, 20],
       [204,  1, 141, 43, 81,  1, 506,  0, 23,  0],
       [ 0,  0,  0,  0,  0, 76,  0, 888,  1, 35],
       [ 8,  1,  3,  4,  2,  9, 48,  4, 921,  0],
       [ 0,  0,  0,  0,  0, 40,  1, 44,  0, 915]])
```

```
#Lets make a better confusion matrix
make_confusion_matrix(y_true=test_labels,
                       y_pred=y_preds,
                       classes=class_names,
                       figsize=(15,15),
                       text_size = 10)
```

Confusion Matrix

True label	T-shirt/top	829 (82.9%)	4 (0.4%)	19 (1.9%)	61 (6.1%)	1 (0.1%)	1 (0.1%)	72 (7.2%)	0 (0.0%)	13 (1.3%)	0
	Trousers	9 (0.9%)	927 (92.7%)	4 (0.4%)	52 (5.2%)	2 (0.2%)	1 (0.1%)	5 (0.5%)	0 (0.0%)	0 (0.0%)	0
	Pullover	29 (2.9%)	1 (0.1%)	743 (74.3%)	18 (1.8%)	105 (10.5%)	0 (0.0%)	103 (10.3%)	0 (0.0%)	1 (0.1%)	0
	Dress	61 (6.1%)	15 (1.5%)	31 (3.1%)	785 (78.5%)	50 (5.0%)	1 (0.1%)	51 (5.1%)	0 (0.0%)	6 (0.6%)	0
	Coat	1 (0.1%)	0 (0.0%)	190 (19.0%)	42 (4.2%)	636 (63.6%)	0 (0.0%)	125 (12.5%)	0 (0.0%)	6 (0.6%)	0
	Sandal	0 (0.0%)	0 (0.0%)	0 (0.0%)	1 (0.1%)	1 (0.1%)	943 (94.3%)	0 (0.0%)	28 (2.8%)	7 (0.7%)	20
	Shirt	204 (20.4%)	1 (0.1%)	141 (14.1%)	43 (4.3%)	81 (8.1%)	1 (0.1%)	506 (50.6%)	0 (0.0%)	23 (2.3%)	0

▼ Let's create a function to plot a random image along with its prediction.

```
import random

# Create a function for plotting a random image along with its prediction
def plot_random_image(model, images, true_labels, classes):
    """Picks a random image, plots it and labels it with a predicted and truth label.

    Args:
        model: a trained model (trained on data similar to what's in images).
        images: a set of random images (in tensor form).
        true_labels: array of ground truth labels for images.
        classes: array of class names for images.

    Returns:
        A plot of a random image from `images` with a predicted class label from `model`
        as well as the truth class label from `true_labels`.
    """
    # Setup random integer
    i = random.randint(0, len(images))

    # Create predictions and targets
    target_image = images[i]
    pred_probs = model.predict(target_image.reshape(1, 28, 28)) # have to reshape to get into right size for model
    pred_label = classes[pred_probs.argmax()]
    true_label = classes[true_labels[i]]

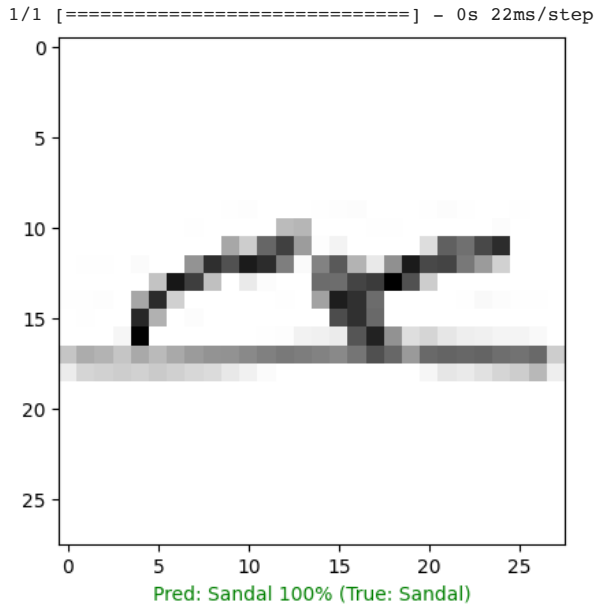
    # Plot the target image
    plt.imshow(target_image, cmap=plt.cm.binary)

    # Change the color of the titles depending on if the prediction is right or wrong
    if pred_label == true_label:
        color = "green"
    else:
        color = "red"

    # Add xlabel information (prediction/truth label)
```

```
plt.xlabel("Pred: {} {:.2f}% (True: {})".format(pred_label,
                                                100*tf.reduce_max(pred_probs),
                                                true_label),
          color=color) # set the color to green or red
```

```
# Check out a random image as well as its prediction
plot_random_image(model=model_4,
                  images=test_data,
                  true_labels=test_labels,
                  classes=class_names)
```



▼ Patterns our model is learning

```
# Find the layers of our most recent model
model_4.layers
```

```
[<keras.layers.resizing.flatten.Flatten at 0x7b90837b0700>,
 <keras.layers.core.dense.Dense at 0x7b9082a0e0b0>,
 <keras.layers.core.dense.Dense at 0x7b9082a0ebc0>,
 <keras.layers.core.dense.Dense at 0x7b9082a0e440>]
```

```
# Extracting a particular layer
model_4.layers[1]
```

```
<keras.layers.core.dense.Dense at 0x7b9082a0e0b0>
```

```
# Getting the patterns of a layer in our network
weights, biases = model_4.layers[1].get_weights()
```

```
# Shape = 1 weight matrix the size of our input data (28x28) per neuron (4)
weights, weights.shape
```

```
(array([[ -0.6286177 ,  0.04304771,  0.27682275, -1.1177504 ],
        [-0.50602835,  1.4612656 , -1.3399148 , -0.69901824],
        [-1.6388192 ,  1.4361247 ,  1.1517581 ,  0.20172885],
        ...,
        [-0.19037662, -0.10284518,  0.09389258, -0.10114335],
        [ 0.52314425,  0.522483 ,  0.27723047, -0.4458822 ],
        [-0.16856955, -0.40255272,  0.59475726,  0.01509975]]),
 dtype=float32),
 (784, 4))
```

The weights matrix is the same shape as the input data, which in our case is 784 (28x28 pixels). And there's a copy of the weights matrix for each neuron in the selected layer (our selected layer has 4 neurons).

Each value in the weights matrix corresponds to how a particular value in the input data influences the network's decisions.

Now let's check out the bias vector.

```
# Shape = 1 bias per neuron (we use 4 neurons in the first layer)
biases, biases.shape

(array([1.6029571, 2.7418947, 4.0676045, 1.9065827], dtype=float32), (4,))
```

Every neuron has a bias vector. Each of these is paired with a weight matrix.

The bias values get initialized as zeroes by default (using the `bias_initializer` parameter).

The bias vector dictates how much the patterns within the corresponding weights matrix should influence the next layer.

```
#Now lets calculate the number of paramters in our model
model_4.summary()
```

Model: "sequential_10"

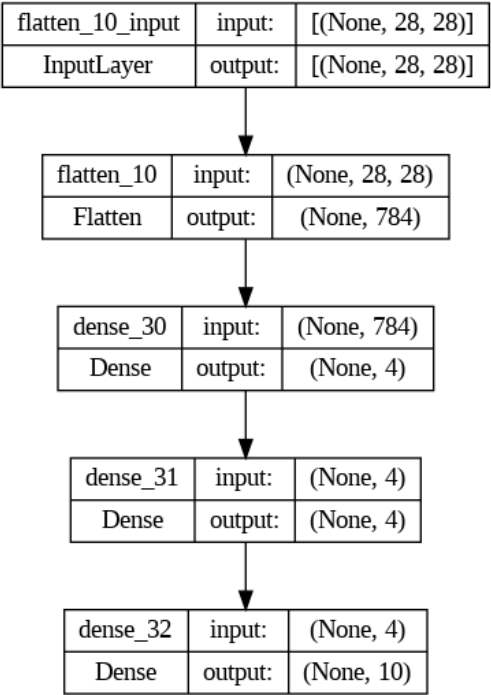
Layer (type)	Output Shape	Param #
flatten_10 (Flatten)	(None, 784)	0
dense_30 (Dense)	(None, 4)	3140
dense_31 (Dense)	(None, 4)	20
dense_32 (Dense)	(None, 10)	50

Total params: 3,210
Trainable params: 3,210
Non-trainable params: 0

Starting from the input layer, each subsequent layer's input is the output of the previous layer as shown below

```
from tensorflow.keras.utils import plot_model

# See the inputs and outputs of each layer
plot_model(model_4, show_shapes=True)
```



```
from google.colab.patches import cv2_imshow
img = cv2.imread('/content/pexels-vie-studio-4439457.jpg', cv2.IMREAD_UNCHANGED)
resized_image = cv2.resize(img, (700,300))
cv2_imshow(resized_image)
```



✓ 2s completed at 1:19 AM

